

# Bios 6301: Assignment 7

Jeongwon Choi

Due Thursday, 03 November, 1:00 PM

$5^{n=day}$  points taken off for each day late.

40 points total.

Submit a single knitr file (named `homework7.rmd`), along with a valid PDF output file. Inside the file, clearly indicate which parts of your responses go with which problems (you may use the original homework document as a template). Add your name as `author` to the file's metadata section. Raw R code/output or word processor files are not acceptable.

Failure to name file `homework7.rmd` or include author name may result in 5 points taken off.

```
library(tidyverse)

## — Attaching packages — tidyverse 1.3.2 —
## ✓ ggplot2 3.3.6      ✓ purrr 0.3.4
## ✓ tibble 3.1.8       ✓ dplyr 1.0.10
## ✓ tidyr 1.2.1        ✓ stringr 1.4.1
## ✓ readr 2.1.2        ✓ forcats 0.5.2
## — Conflicts — tidyverse_conflicts() —
## * dplyr::filter() masks stats::filter()
## * dplyr::lag()     masks stats::lag()
```

## Question 1

21 points

Use the following code to generate data for patients with repeated measures of A1C (a test for levels of blood glucose).

```
genData <- function(n) {
  if(exists(".Random.seed", envir = .GlobalEnv)) {
    save.seed <- get(".Random.seed", envir= .GlobalEnv)
    on.exit(assign(".Random.seed", save.seed, envir = .GlobalEnv))
  } else {
    on.exit(rm(".Random.seed", envir = .GlobalEnv))
  }
  set.seed(n)
  subj <- ceiling(n / 10)
  id <- sample(subj, n, replace=TRUE)
  times <- as.integer(difftime(as.POSIXct("2005-01-01"), as.POSIXct("2000-01-01"), units='secs'))
  dt <- as.POSIXct(sample(times, n), origin='2000-01-01')
  mu <- runif(subj, 4, 10)
  alc <- unsplit(mapply(rnorm, tabulate(id), mu, SIMPLIFY=FALSE), id)
  data.frame(id, dt, alc)
}
x <- genData(500)
```

Perform the following manipulations: (3 points each)

- 1. Order the data set by `id` and `dt`.

```
x <- x[order(x$id, x$dt),]
rownames(x) <- NULL
head(x)
```

	id <int>	dt <dtm>	a1c <dbl>
1	1	2001-05-08 16:22:52	7.309995
2	1	2001-06-17 22:42:23	8.310721
3	1	2001-08-17 16:51:46	6.548845
4	1	2001-12-14 14:50:29	5.985275



```

nr <- nrow(x)

for (i in 1:nr){
  this_idx <- which(x$id==i)
  this_mark <- x[this_idx, "mark"]
  if(this_mark==0 ){
    next} else{
  for (j in 1:this_mark){

    new.dt <- x[this_idx-1, 'dt'] + as.difftime(365, units="days")
    x <- x %>% add_row(id=x[this_idx, "id"], dt=new.dt, a1c=NA, .before=this_idx)
    # x[this_idx-1, 'gap'] <- x$dt[this_idx-1] - x$dt[this_idx-2]
    x$gap[[this_idx]] = x$dt[this_idx] - x$dt[this_idx-1]
  }
}
}
x

```

id <int>	dt <dtm>	a1c <dbl>	gap <dbl>	mark <dbl>	idx <int>
1	2001-05-08 16:22:52	7.309995	0.000000	0	1
1	2001-06-17 22:42:23	8.310721	40.263553	0	2
1	2001-08-17 16:51:46	6.548845	60.756516	0	3
1	2001-12-14 14:50:29	5.985275	118.957442	0	4
1	2002-08-19 13:51:47	6.011547	247.917569	0	5
1	2003-03-22 03:51:36	7.243858	214.624873	0	6
1	2003-06-27 01:01:34	5.170870	96.840255	0	7
2	2001-03-05 22:24:43	9.237660	0.000000	0	8
2	2001-03-16 17:45:49	11.637444	10.806319	0	9
2	2001-05-02 04:14:56	10.085473	46.395220	0	10

1-10 of 545 rows

Previous 1 2 3 4 5 6 ... 55 Next

3. Create a new column `visit`. For each `id`, add the visit number. This should be 1 to `n` where `n` is the number of observations for an individual. This should include the observations created with missing `a1c` values.

```

library(tidyverse)
x$visit <- NULL
uniqueid <- unique(x$id)
for (i in uniqueid) {
  this_id <- which(x$id==i)
  count_this_id <- length(this_id)
  x[this_id, 'visit'] <- 1:count_this_id
}

```

4. For each `id`, replace missing values with the mean `a1c` value for that individual.

```

#create a mean a1c value
mean_a1c <- x %>%
group_by(id) %>%
summarize(Mean = mean(a1c, na.rm=TRUE))
mean_a1c

```

id <int>	Mean <dbl>
1	6.654444
2	9.789132
3	6.951820

id <int>	Mean <dbl>
4	8.191985
5	9.429694
6	7.133443
7	7.879138
8	6.244061
9	4.420523
10	6.028370

1-10 of 50 rows

Previous 1 2 3 4 5 Next

```
#replace NA with the mean
uniqueid2 <- unique(x$id)
missingidx <- is.na(x$alc)

for(i in uniqueid2) {
  this_id2 <- x$id==i
  both <- this_id2 & missingidx
  x[both,'alc'] <- mean_alc[i,'Mean']
}
```

5. Print mean alc for each id.

```
x %>%
group_by(id) %>%
summarize(Mean = mean(alc))
```

id <int>	Mean <dbl>
1	6.654444
2	9.789132
3	6.951820
4	8.191985
5	9.429694
6	7.133443
7	7.879138
8	6.244061
9	4.420523
10	6.028370

1-10 of 50 rows

Previous 1 2 3 4 5 Next

```
mean_alc
```

id <int>	Mean <dbl>
1	6.654444
2	9.789132
3	6.951820
4	8.191985
5	9.429694

<b>id</b> <int>	<b>Mean</b> <dbl>
6	7.133443
7	7.879138
8	6.244061
9	4.420523
10	6.028370

1-10 of 50 rows

Previous **1** 2 3 4 5 Next

6. Print total number of visits for each `id` .

```
lastobs <- x %>%
  group_by(id) %>%
  summarise_all(last)

lastobs[,c(1,7)]
```

<b>id</b> <int>	<b>visit</b> <int>
1	7
2	16
3	13
4	9
5	14
6	11
7	7
8	12
9	15
10	8

1-10 of 50 rows

Previous **1** 2 3 4 5 Next

7. Print the observations for `id = 15` .

```
x[x$id == 15,]
```

	<b>id</b> <int>	<b>dt</b> <dtm>	<b>a1c</b> <dbl>	<b>gap</b> <dbl>	<b>mark</b> <dbl>	<b>idx</b> <int>	<b>visit</b> <int>
158	15	2000-10-21 01:08:17	7.401322	0.000000	0	144	1
159	15	2001-08-08 14:23:08	5.896318	291.551979	0	145	2
160	15	2001-08-15 07:03:29	7.457722	6.694687	0	146	3
161	15	2002-03-15 21:23:10	5.330917	212.638669	0	147	4
162	15	2002-04-14 09:08:25	6.484003	29.448090	0	148	5
163	15	2002-10-10 18:27:43	8.139101	179.388403	0	149	6
164	15	2003-02-19 12:58:53	6.446557	131.813310	0	150	7
165	15	2003-03-02 06:58:10	7.432291	10.749502	0	151	8
166	15	2003-06-30 07:20:49	7.113792	119.974063	0	152	9
167	15	2004-01-22 20:30:42	5.668897	206.590197	0	153	10

1-10 of 10 rows

## Question 2

16 points

Install the `lexicon` package. Load the `sw_fry_1000` vector, which contains 1,000 common words.

```
library(lexicon)
data('sw_fry_1000', package = 'lexicon')
head(sw_fry_1000)
```

```
## [1] "the" "of" "to" "and" "a" "in"
```

1. Remove all non-alphabetical characters and make all characters lowercase. Save the result as `a`.

```
a1 <- tolower(sw_fry_1000)
a <- gsub("[^a-z]", "", a1)
```

Use vector `a` for the following questions. (2 points each)

2. How many words contain the string “ar”?

```
length(grep("ar", a, value=TRUE))
```

```
## [1] 64
```

64 words contain the string “ar”.

3. Find a six-letter word that starts with “l” and ends with “r”.

```
b1 <- grep("^l", a, value=TRUE)
b2 <- grep("r$", b1, value=TRUE)
str_length(b2)
```

```
## [1] 6
```

letter starts with “l” and ends with “r”.

4. Return all words that start with “col” or end with “eck”.

```
grep("^col|eck$", a, value=TRUE)
```

```
## [1] "color" "cold" "check" "collect" "colony" "column" "neck"
```

5. Find the number of words that contain 4 or more adjacent consonants. Assume “y” is always a consonant.

```
length(grep("[^aeiou]{4}", a, value=TRUE))
```

```
## [1] 8
```

8 words contain 4 or more adjacent consonants

6. Return all words with a “q” that isn’t followed by a “ui”.

```
c1 <- grep("[q]", a, value=TRUE)
idx <- grep("[q](?!ui)", c1)
c1[-idx]
```

```
## [1] "question" "equate" "square" "equal" "quart" "quotient"
```

7. Find all words that contain a “k” followed by another letter. Run the `table` command on the first character following the first “k” of each word.

```
# I first find words that contain a "k" followed by another letter
d <- grep("[k].+[[:alpha:]]", a, value=TRUE)
# Then split words into a vector of letters
d1 <- strsplit(d, '')
# For vector d1, find where k is located within a word
d2 <- unlist(lapply(d1, function(x){y= x=='k'; return(which(y))}))

# Make a loop with the location of k, return the letter followed by k.
d3 <- c()
for (i in 1:length(d2)){
  d3<- c( d3,d1[[i]][d2[i]+1])
}

# Make a table with d3.
table(d3)
```

```
## d3
## e i n
## 4 5 2
```

8. Remove all vowels. How many character strings are found exactly once?

```
e1 <- gsub("[aeiou]", '', a)
sum(table(e1)==1)
```

```
## [1] 581
```

581 character strings are found once.

## Question 3

3 points

The first argument to most functions that fit linear models are formulas. The following example defines the response variable `death` and allows the model to incorporate all other variables as terms. `.` is used to mean all columns not otherwise in the formula. I

```
# I changed the dataset because error message came out
haart_df <- read.csv('~/downloads/haart.csv')[,c('death','weight','hemoglobin','cd4baseline')]
coef(summary(glm(death ~ ., data=haart_df, family=binomial(logit))))
```

```
##              Estimate Std. Error   z value    Pr(>|z|)
## (Intercept)  3.576411744 1.226870535   2.915069 0.0035561039
## weight      -0.046210552 0.022556001  -2.048703 0.0404911395
## hemoglobin  -0.350642786 0.105064078  -3.337418 0.0008456055
## cd4baseline  0.002092582 0.001811959   1.154872 0.2481427160
```

Now imagine running the above several times, but with a different response and data set each time. Here's a function:

```
myfun <- function(dat, response) {
  form <- as.formula(response ~ .)
  coef(summary(glm(form, data=dat, family=binomial(logit))))
}
```

Unfortunately, it doesn't work. `tryCatch` is "catching" the error so that this file can be knit to PDF.

```
tryCatch(myfun(haart_df, death), error = function(e) e)
```

```
## <simpleError in eval(predvars, data, env): object 'death' not found>
```

What do you think is going on? Consider using `debug` to trace the problem.

The 'response' argument in the `myfun` function is not recognized as an object. So it doesn't work because R cannot recognize the variable such as `death` or `hemoglobin`, which makes it impossible to form a new function.

5 bonus points

Create a working function.

```
myfun_new <- function(dat, response) {  
  form <- as.formula( paste0(substitute(response), "~.") )  
  coef(summary(glm(form, data=dat, family=binomial(logit))))  
}  
  
myfun_new(haart_df, death)
```

```
##              Estimate Std. Error   z value    Pr(>|z|)  
## (Intercept)  3.576411744  1.226870535   2.915069 0.0035561039  
## weight      -0.046210552  0.022556001  -2.048703 0.0404911395  
## hemoglobin  -0.350642786  0.105064078  -3.337418 0.0008456055  
## cd4baseline  0.002092582  0.001811959   1.154872 0.2481427160
```

So in this case, I included `substitute(response)` so that the function could recognize the response variable name as a symbol for formula and use that symbol to create a formula.